

In [6]:

```
#Necessary import
import gym
import random
import numpy as np
import tensorflow.compat.v1 as tf
tf.disable_v2_behavior()
from collections import deque
```

In [7]:

```
#CartPole environment import
env_name = "CartPole-v0"
env = gym.make(env_name)

#Show moves available and space dimensions
print("Observation space:", env.observation_space)
print("Action space:", env.action_space)
```

```
Observation space: Box([-4.8000002e+00 -3.4028235e+38 -4.1887903e-01 -3.4028
235e+38], [4.8000002e+00 3.4028235e+38 4.1887903e-01 3.4028235e+38], (4,), f
loat32)
Action space: Discrete(2)
```

In [8]:

```
#Neural network definition
class QNetwork():
    def __init__(self, state_dim, action_size):
        #Create tensorflow variable for environment characteristics
        self.state_in = tf.placeholder(tf.float32, shape=[None, *state_dim])
        self.action_in = tf.placeholder(tf.int32, shape=[None])
        self.q_target_in = tf.placeholder(tf.float32, shape=[None])
        action_one_hot = tf.one_hot(self.action_in, depth=action_size)

        #Create layers of the network
        self.hidden1 = tf.layers.dense(self.state_in, 100, activation=tf.nn.relu)
        self.q_state = tf.layers.dense(self.hidden1, action_size, activation=None)
        self.q_state_action = tf.reduce_sum(tf.multiply(self.q_state, action_one_hot), axis=1)

        #Loss function instance
        self.loss = tf.reduce_mean(tf.square(self.q_state_action - self.q_target_in))

        #Create optimizer instance
        self.optimizer = tf.train.AdamOptimizer(learning_rate=0.001).minimize(self.loss)

    #Model update
    def update_model(self, session, state, action, q_target):
        feed = {self.state_in: state, self.action_in: action, self.q_target_in: q_target}
        session.run(self.optimizer, feed_dict=feed)

    #Run the network
    def get_q_state(self, session, state):
        q_state = session.run(self.q_state, feed_dict={self.state_in: state})
        return q_state
```

In [9]:

```

class ReplayBuffer():
    def __init__(self, maxlen):
        self.buffer = deque(maxlen=maxlen)

    def add(self, experience):
        self.buffer.append(experience)

    def sample(self, batch_size):
        sample_size = min(len(self.buffer), batch_size)
        samples = random.choices(self.buffer, k=sample_size)
        return map(list, zip(*samples))

```

In [10]:

```

#Deep Q-Learning agent
class DQNAgent():
    def __init__(self, env):
        #Instance of the environment characteristics
        self.state_dim = env.observation_space.shape
        self.action_size = env.action_space.n

        #Instance of Q-Learning Network
        self.q_network = QNetwork(self.state_dim, self.action_size)

        #Instance of the replay buffer
        self.replay_buffer = ReplayBuffer(maxlen=10000)

        #Costant of the agent
        self.gamma = 0.99
        self.eps = 1.0

        #Create and run session
        self.sess = tf.Session()
        self.sess.run(tf.global_variables_initializer())

    #Function that take the action to perform to the environment
    def get_action(self, state):
        q_state = self.q_network.get_q_state(self.sess, [state])
        action_greedy = np.argmax(q_state)
        action_random = np.random.randint(self.action_size)
        action = action_random if random.random() < self.eps else action_greedy
        return action

    #Function that want to update the agent after one action performed
    def train(self, state, action, next_state, reward, done):
        self.replay_buffer.add((state, action, next_state, reward, done))
        states, actions, next_states, rewards, dones = self.replay_buffer.sample(50)
        q_next_states = self.q_network.get_q_state(self.sess, next_states)
        q_next_states[dones] = np.zeros([self.action_size])
        q_targets = rewards + self.gamma * np.max(q_next_states, axis=1)
        self.q_network.update_model(self.sess, states, actions, q_targets)

        if done: self.eps = max(0.1, 0.99*self.eps)

    def __del__(self):
        self.sess.close()

```

In [11]:

```

#Agent instance
agent = DQNAgent(env)

#Episodes number
num_episodes = 400
total_reward_array=[]

#Each iteration represent one episode
for ep in range(num_episodes):
    #Reset environment
    state = env.reset()

    total_reward = 0
    done = False
    while not done:
        #Take action from agent and permorm it
        action = agent.get_action(state)
        next_state, reward, done, info = env.step(action)

        #Update agent
        agent.train(state, action, next_state, reward, done)

        #Show environment state
        env.render()

        total_reward += reward
        state = next_state

    #Show episode total reward
    total_reward_array.append(total_reward)
    print("Episode: {}, total_reward: {:.2f}".format(ep, total_reward))

```

```

Episode: 354, total_reward: 200.00
Episode: 355, total_reward: 200.00
Episode: 356, total_reward: 200.00
Episode: 357, total_reward: 200.00
Episode: 358, total_reward: 200.00
Episode: 359, total_reward: 200.00
Episode: 360, total_reward: 200.00
Episode: 361, total_reward: 200.00

Episode: 362, total_reward: 161.00
Episode: 363, total_reward: 200.00
Episode: 364, total_reward: 200.00
Episode: 365, total_reward: 200.00
Episode: 366, total_reward: 200.00
Episode: 367, total_reward: 200.00
Episode: 368, total_reward: 200.00
Episode: 369, total_reward: 200.00
Episode: 370, total_reward: 200.00
Episode: 371, total_reward: 200.00
Episode: 372, total_reward: 189.00
Episode: 373, total_reward: 160.00

```

In [12]:

```

#Close environment
env.close()

```

In [13]:

```
#Necessary import to plot total reward results  
import seaborn as sns  
import matplotlib.pyplot as plt
```

In [14]:

```
#Plot total reward results  
fig, ax = plt.subplots(figsize=(16,8))  
sns.lineplot(x=range(0,400), y= total_reward_array, ax=ax)  
plt.title('Cartpole DQN agent')  
plt.show
```

Out[14]:

```
<function matplotlib.pyplot.show(close=None, block=None)>
```

